

Chapter 4: Threads & Concurrency



Chapter 4: Threads

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples



Objectives

- ❑ Identify the *basic components of a thread*, and *contrast threads and processes*
- ❑ Describe the *benefits* and *challenges* of designing *multithreaded applications*
- ❑ Illustrate *different approaches to implicit threading* including thread pools, fork-join, and Grand Central Dispatch
- ❑ Describe how the Windows and Linux operating systems represent threads
- ❑ Design multithreaded applications using the Pthreads, Java, and Windows threading APIs

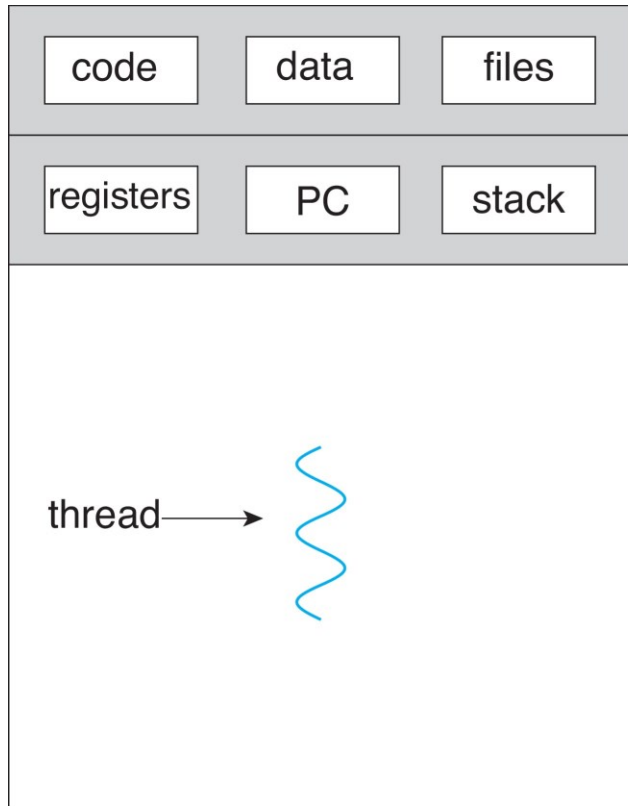


Motivation


- ❑ Most modern applications are *multithreaded*
- ❑ *Threads run within application*
- ❑ Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- ❑ Process creation is heavy-weight while *thread creation is light-weight*
- ❑ Can simplify code, increase efficiency
- ❑ Kernels are generally multithreaded



Single and Multithreaded Processes

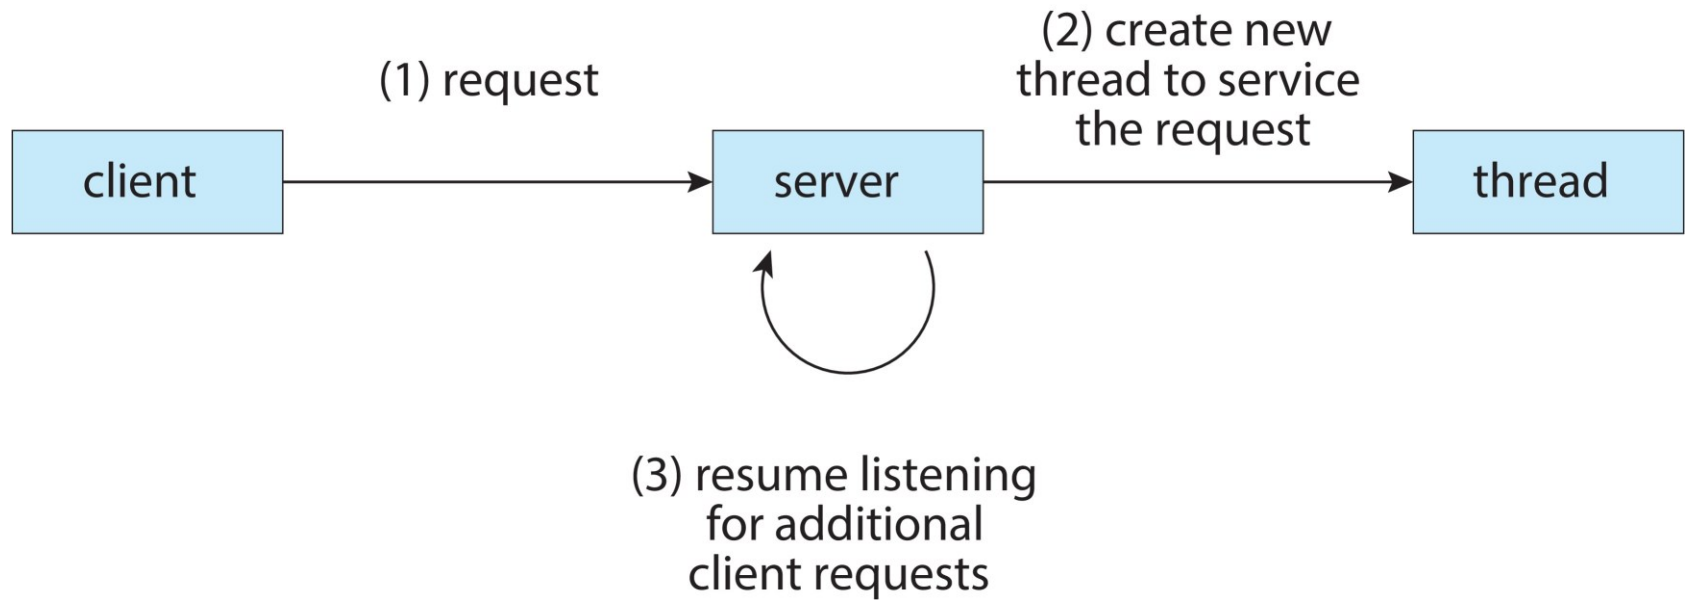


single-threaded process



multithreaded process

Multithreaded Server Architecture



Benefits

- ❑ *Responsiveness* – may allow continued execution if part of process is blocked, especially important for user interfaces
- ❑ *Resource Sharing* – threads share resources of process, easier than shared memory or message passing (IPC)
- ❑ *Economy* – cheaper than process creation, thread switching lower overhead than context switching
- ❑ *Scalability* – process can take advantage of multicore architectures




Multicore Programming

- ❑ *Multicore* or *multiprocessor systems* putting pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- ❑ *Parallelism* implies a system can perform more than one task simultaneously
- ❑ *Concurrency* supports more than one task making progress
 - Single processor / core, scheduler providing concurrency




Concurrency vs. Parallelism

❑ Concurrent execution on single-core system:



❑ Parallelism on a multi-core system:



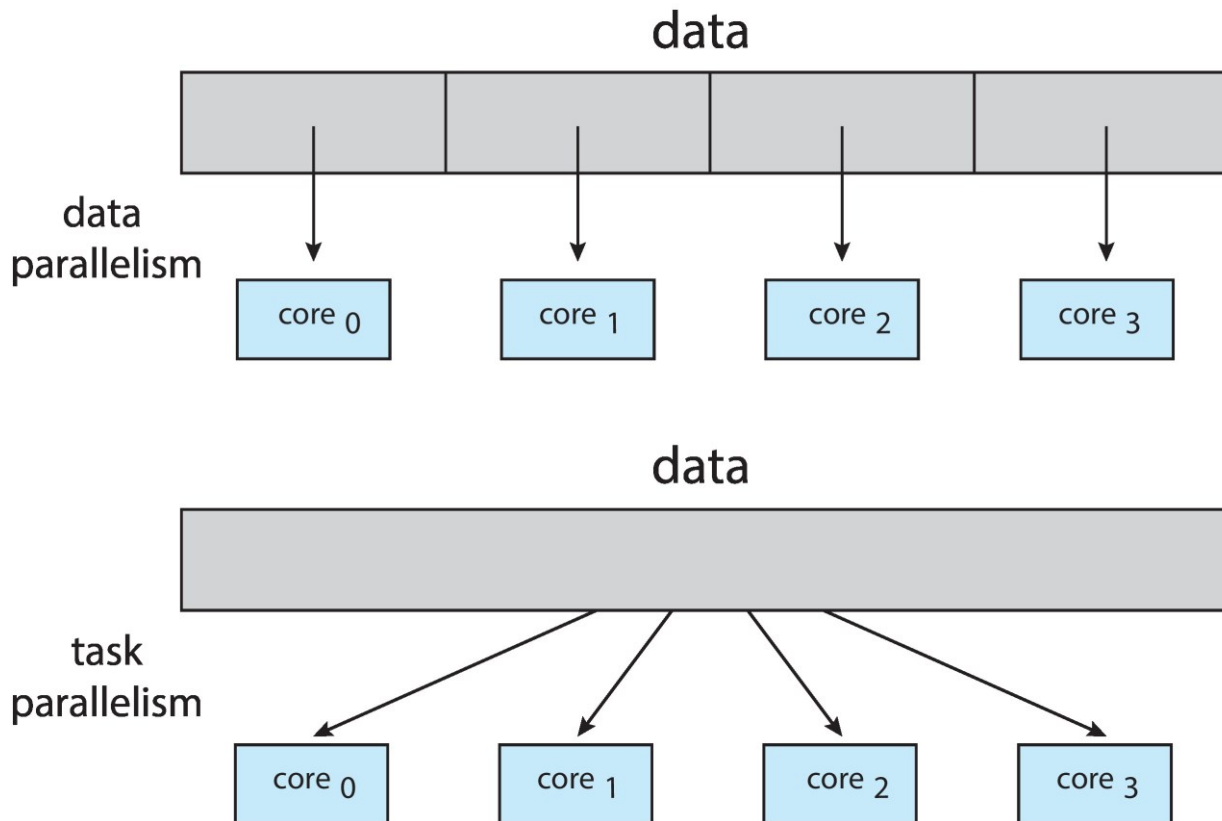
Multicore Programming

□ Types of parallelism

- **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
- **Task parallelism** – distributing threads across cores, each thread performing unique operation



Data and Task Parallelism




Amdahl's Law

- ❑ Identifies *performance gains* from adding additional cores to an application that has both serial and parallel components
 - **S** is serial portion
 - **N** processing cores
 - That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
 - As **N** approaches infinity, speedup approaches $1/S$
- ❑ Serial portion of an application has disproportionate effect on performance gained by adding additional cores
- ❑ But does the law take into account *contemporary multicore systems*?

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$



Amdahl's Law




User Threads and Kernel Threads

- ❑ *User threads* - management done by user-level threads library
- ❑ Three primary thread libraries:
 - POSIX Pthreads
 - Windows threads
 - Java threads
- ❑ *Kernel threads* - supported by the Kernel
- ❑ Examples – virtually all general purpose operating systems, including:
 - Windows, Linux, Mac OS X
 - iOS, Android



User and Kernel Threads




Multithreading Models

- ☐ Many-to-One
- ☐ One-to-One
- ☐ Many-to-Many



Many-to-One


- ❑ *Many user-level threads* mapped to *single kernel thread*
- ❑ One thread blocking causes all to block
- ❑ Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- ❑ *Few systems currently use this model*
- ❑ Examples:
 - Solaris Green Threads
 - GNU Portable Threads



One-to-One


- ❑ *Each user-level thread* maps to *one kernel thread*
- ❑ Creating a user-level thread creates a kernel thread
- ❑ More concurrency than many-to-one
- ❑ Number of threads per process sometimes restricted due to overhead
- ❑ Examples

- Windows
- Linux



Many-to-Many Model

- ❑ Allows *many user level threads* to be mapped to *many kernel threads*
- ❑ Allows the operating system to create a sufficient number of kernel threads
- ❑ Windows with the **ThreadFiber** package
- ❑ Otherwise *not very common*



Thread Libraries

- ❑ **Thread library** provides programmer with API for creating and managing threads
- ❑ Two primary ways of implementing
 - Library entirely *in user space*
 - Kernel-level library *supported by the OS*



Pthreads

- ❑ May be provided either as *user-level* or *kernel-level*
- ❑ A **POSIX standard (IEEE 1003.1c) API** for thread creation and synchronization
- ❑ Specification, not implementation
- ❑ API specifies behavior of the thread library, implementation is up to development of the library
- ❑ Common in UNIX operating systems (Linux & Mac OS X)



Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```



Pthreads Example (cont)

```
/* The thread will execute in this function */  
void *runner(void *param)  
{  
    int i, upper = atoi(param);  
    sum = 0;  
  
    for (i = 1; i <= upper; i++)  
        sum += i;  
  
    pthread_exit(0);  
}
```



Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



End of Chapter 4

